


dition if H were a *complete* history. Otherwise there are problems. First, there is an artificial problem in that a partial history can never be equivalent to a serial one. This is because serial histories must be complete by definition, and two histories can be equivalent only if they contain the same set of operations. But even assuming for the moment that we did not insist on serial histories being complete, we would still have a problem—a problem of substance and not mere definition. Namely, an incomplete execution of a transaction does not necessarily preserve the consistency of the database. Thus, serializability would be an inappropriate correctness condition if it merely stated that a history be computationally equivalent to a serial execution of some possibly partial transactions, since such a history does not necessarily represent a consistency preserving execution.

We are therefore led to a slightly more complex definition of serializability. Although more complex, the definition is still natural and, most importantly, sound. A history H is *serializable* (SR) if its committed projection, $C(H)$, is equivalent to a serial history H_s .

This takes care of the problems, because $C(H)$ is a *complete* history. Moreover, it is not an arbitrarily chosen complete history. If H represents the execution so far, it is really only the *committed* transactions whose execution the DBS has unconditionally guaranteed. All other transactions may be aborted.

2.3 THE SERIALIZABILITY THEOREM

We can determine whether a history is serializable by analyzing a graph derived from the history called a serialization graph. Let H be a history over $T = \{T_1, \dots, T_n\}$. The *serialization graph* (SG) for H , denoted $SG(H)$, is a directed graph whose nodes are the transactions in T that are committed in H and whose edges are all $T_i \rightarrow T_j$ ($i \neq j$) such that one of T_i 's operations precedes and conflicts with one of T_j 's operations in H . For example:



The edge $T_1 \rightarrow T_3$ is in $SG(H_s)$ because $w_1[x] < r_3[x]$, and the edge $T_2 \rightarrow T_3$ is in $SG(H_s)$ because $r_2[x] < w_3[x]$. Notice that a single edge in $SG(H_s)$ can be present because of more than one pair of conflicting operations. For instance, the edge $T_2 \rightarrow T_1$ is caused both by $r_2[x] < w_1[x]$ and $w_2[y] < w_1[y]$.

In general, the existence of edges $T_i \rightarrow T_j$ and $T_j \rightarrow T_k$ in an SG does not necessarily imply the existence of edge $T_i \rightarrow T_k$. For instance, with $w_3[z]$ replacing $w_3[x]$ in T_3 , $SG(H_s)$ becomes

$$T_2 \rightarrow T_1 \rightarrow T_3$$

since there is no conflict between T_2 and T_3 .⁷

Each edge $T_i \rightarrow T_j$ in $SG(H)$ means that at least one of T_i 's operations precedes and conflicts with one of T_j 's. This suggests that T_i should precede T_j in any serial history that is equivalent to H . If we can find a serial history, H_s , consistent with all edges in $SG(H)$, then $H_s \equiv H$ and so H is SR. We can do this as long as $SG(H)$ is acyclic.

In our previous example, $SG(H_s)$ is acyclic. A serial history where transactions appear in an order consistent with the edges of $SG(H_s)$ is $T_2 T_1 T_3$. Indeed this is the only such serial history. You can easily verify that H_s is equivalent to $T_2 T_1 T_3$ and is therefore SR. We formalize this intuitive argument in the following theorem—the fundamental theorem of serializability theory.

Theorem 2.1: (The Serializability Theorem) A history H is serializable iff $SG(H)$ is acyclic.

Proof: (if) Suppose H is a history over $T = \{T_1, T_2, \dots, T_n\}$. Without loss of generality, assume T_1, T_2, \dots, T_m ($m \leq n$) are all of the transactions in T that are committed in H . Thus T_1, T_2, \dots, T_m are the nodes of $SG(H)$. Since $SG(H)$ is acyclic it may be topologically sorted.⁸ Let i_1, \dots, i_m be a permutation of $1, 2, \dots, m$ such that $T_{i_1}, T_{i_2}, \dots, T_{i_m}$ is a topological sort of $SG(H)$. Let H_s be the serial history $T_{i_1} T_{i_2} \dots T_{i_m}$. We claim that $C(H) \equiv H_s$. To see this, let $p_i \in T_i$ and $q_j \in T_j$, where T_i, T_j are committed in H . Suppose p_i, q_j conflict and $p_i <_H q_j$. By definition of $SG(H)$, $T_i \rightarrow T_j$ is an edge in $SG(H)$. Therefore in any topological sort of $SG(H)$, T_i must appear before T_j . Thus in H_s all operations of T_i appear before any operation of T_j , and in particular, $p_i <_{H_s} q_j$. We have proved that any two conflicting operations are ordered in $C(H)$ in the same way as H_s . Thus $C(H) \equiv H_s$ and, because H_s is serial by construction, H is SR as was to be proved.

(only if) Suppose history H is SR. Let H_s be a serial history equivalent to $C(H)$. Consider an edge $T_i \rightarrow T_j$ in $SG(H)$. Thus there are two conflicting operations p_i, q_j of T_i, T_j (respectively), such that $p_i <_H q_j$. Because $C(H) \equiv H_s$, $p_i <_{H_s} q_j$. Because H_s is serial and p_i in T_i precedes q_j in T_j , it follows that T_i appears before T_j in H_s . Thus, we've shown that if $T_i \rightarrow T_j$ is in $SG(H)$ then T_i appears before T_j in H_s . Now suppose there is a cycle in $SG(H)$, and without loss of generality let that cycle be $T_1 \rightarrow T_2 \rightarrow \dots \rightarrow T_k \rightarrow T_1$. These edges imply that in H_s , T_1 appears before T_2 which appears before T_3 which appears ... before T_k which appears before T_1 . Thus, the existence of the cycle implies that each of T_1, T_2, \dots, T_k appears before

⁷We say that *two transactions conflict* if they contain conflicting operations.

⁸See Section A.3 of the Appendix for a definition of “topological sort of a directed acyclic graph.”

itself in the serial history H_s , an absurdity. So no cycle can exist in $SG(H)$. That is, $SG(H)$ is an acyclic directed graph, as was to be proved. \square

From the proof of the “if” part of this theorem we see that if a complete history H has an acyclic $SG(H)$, then H is equivalent to *any* serial history that’s a topological sort of $SG(H)$. Since the latter can have more than one topological sort, H may be equivalent to more than one serial history. For instance,

$$H_6 = w_1[x] \ w_1[y] \ c_1 \ r_2[x] \ r_3[y] \ w_2[x] \ c_2 \ w_3[y] \ c_3$$

has the serialization graph

$$SG(H_6) = \begin{array}{c} \xrightarrow{\quad} \\ T_1 \rightarrow T_3 \rightarrow T_2 \end{array}$$

which has two topological sorts, T_1, T_2, T_3 and T_1, T_3, T_2 . Thus H_6 is equivalent to both $T_1 \ T_2 \ T_3$ and $T_1 \ T_3 \ T_2$.

2.4 RECOVERABLE HISTORIES

In Chapter 1, we saw that to ensure correctness in the presence of failures the scheduler must produce executions that are not only SR but also recoverable. We also discussed some additional requirements that may be desirable—preventing cascading aborts and the loss of before images—leading us to the idea of strict executions. Like serializability, these concepts can be conveniently formulated in the language of histories.

A transaction T_i reads data item x from T_j if T_j was the transaction that had last written into x but had not aborted at the time T_i read x . More precisely, we say that T_i reads x from T_j in history H if

1. $w_j[x] < r_i[x]$;
2. $a_j \not< r_i[x]^*$ and
3. if there is some $w_k[x]$ such that $w_j[x] < w_k[x] < r_i[x]$, then $a_k < r_i[x]$.

We say that T_i reads from T_j in H if T_i reads some data item from T_j in H . Notice that it is possible for a transaction to read a data item from itself (i.e., $w_i[x] < r_i[x]$).

A history H is called *recoverable* (RC) if, whenever T_i reads from T_j ($i \neq j$) in H and $c_i \in H, c_j < c_i$. Intuitively, a history is recoverable if each transaction commits after the commitment of all transactions (other than itself) from which it read.

A history H *avoids cascading aborts* (ACA) if, whenever T_i reads x from T_j ($i \neq j$), $c_j < r_i[x]$. That is, a transaction may read only those values that are written by committed transactions or by itself.

* $p \not< q$ denotes that operation p does not precede q in the partial order.